

Liang_Assign5

June 30, 2024

```
[78]: import pandas as pd
import numpy as np
import matplotlib.pyplot as plt
import seaborn as sns; sns.set()
from sklearn.model_selection import StratifiedKFold
from sklearn.metrics import accuracy_score
import warnings
warnings.filterwarnings("ignore")
```

```
[31]: df = pd.read_csv("heart_dataset-1.csv")
```

```
[32]: df.shape
```

```
[32]: (918, 12)
```

```
[33]: df.head()
```

```
[33]:
```

	Age	Sex	ChestPainType	RestingBP	Cholesterol	FastingBS	RestingECG	MaxHR	\
0	40	M	ATA	140	289	0	Normal	172	
1	49	F	NAP	160	180	0	Normal	156	
2	37	M	ATA	130	283	0	ST	98	
3	48	F	ASY	138	214	0	Normal	108	
4	54	M	NAP	150	195	0	Normal	122	

	ExerciseAngina	Oldpeak	ST_Slope	HeartDisease
0	N	0.0	Up	0
1	N	1.0	Flat	1
2	N	0.0	Up	0
3	Y	1.5	Flat	1
4	N	0.0	Up	0

```
[34]: df.columns.values
```

```
[34]: array(['Age', 'Sex', 'ChestPainType', 'RestingBP', 'Cholesterol',
       'FastingBS', 'RestingECG', 'MaxHR', 'ExerciseAngina', 'Oldpeak',
       'ST_Slope', 'HeartDisease'], dtype=object)
```

```
[35]: df.info()
```

```
<class 'pandas.core.frame.DataFrame'>
RangeIndex: 918 entries, 0 to 917
Data columns (total 12 columns):
#   Column                Non-Null Count  Dtype
---  -
0   Age                   918 non-null   int64
1   Sex                   918 non-null   object
2   ChestPainType         918 non-null   object
3   RestingBP             918 non-null   int64
4   Cholesterol            918 non-null   int64
5   FastingBS             918 non-null   int64
6   RestingECG            918 non-null   object
7   MaxHR                 918 non-null   int64
8   ExerciseAngina        918 non-null   object
9   Oldpeak               918 non-null   float64
10  ST_Slope              918 non-null   object
11  HeartDisease          918 non-null   int64
dtypes: float64(1), int64(6), object(5)
memory usage: 86.2+ KB
```

```
[36]: for i in [1,2,6,8,10]:
      print(df.iloc[:,i].unique())
```

```
['M' 'F']
['ATA' 'NAP' 'ASY' 'TA']
['Normal' 'ST' 'LVH']
['N' 'Y']
['Up' 'Flat' 'Down']
```

```
[37]: df_encoded = pd.get_dummies(df, columns=["Sex", "ChestPainType", "RestingECG",
      ↪ "ExerciseAngina", "ST_Slope"])
```

```
df_encoded
```

```
[37]:   Age  RestingBP  Cholesterol  FastingBS  MaxHR  Oldpeak  HeartDisease  \
0    40         140          289          0    172        0.0            0
1    49         160          180          0    156        1.0            1
2    37         130          283          0     98        0.0            0
3    48         138          214          0    108        1.5            1
4    54         150          195          0    122        0.0            0
..   ...         ...         ...         ...   ...   ...
913  45         110          264          0    132        1.2            1
914  68         144          193          1    141        3.4            1
915  57         130          131          0    115        1.2            1
916  57         130          236          0    174        0.0            1
917  38         138          175          0    173        0.0            0

      Sex_F  Sex_M  ChestPainType_ASY  ...  ChestPainType_NAP  \
```

0	False	True	False	...	False
1	True	False	False	...	True
2	False	True	False	...	False
3	True	False	True	...	False
4	False	True	False	...	True
..
913	False	True	False	...	False
914	False	True	True	...	False
915	False	True	True	...	False
916	True	False	False	...	False
917	False	True	False	...	True

	ChestPainType_TA	RestingECG_LVH	RestingECG_Normal	RestingECG_ST	\
0	False	False	True	False	
1	False	False	True	False	
2	False	False	False	True	
3	False	False	True	False	
4	False	False	True	False	
..	
913	True	False	True	False	
914	False	False	True	False	
915	False	False	True	False	
916	False	True	False	False	
917	False	False	True	False	

	ExerciseAngina_N	ExerciseAngina_Y	ST_Slope_Down	ST_Slope_Flat	\
0	True	False	False	False	
1	True	False	False	True	
2	True	False	False	False	
3	False	True	False	True	
4	True	False	False	False	
..	
913	True	False	False	True	
914	True	False	False	True	
915	False	True	False	True	
916	True	False	False	True	
917	True	False	False	False	

	ST_Slope_Up
0	True
1	False
2	True
3	False
4	True
..	...
913	False
914	False

```
915         False
916         False
917         True
```

```
[918 rows x 21 columns]
```

```
[59]: from sklearn.model_selection import cross_val_score
from sklearn.naive_bayes import GaussianNB
from sklearn.svm import LinearSVC
from sklearn.neural_network import MLPClassifier
from sklearn.tree import DecisionTreeClassifier
from sklearn.ensemble import RandomForestClassifier
```

0.1 Problem 1

```
[56]: dfX = df_encoded.loc[:, df_encoded.columns != 'HeartDisease']
dfy = df_encoded.loc[:, df_encoded.columns == 'HeartDisease'].values.ravel()
X = dfX.values
y = dfy

def eval_classifier(_clf, _X, _y):
    accuracies = []
    kf = StratifiedKFold(n_splits=10, shuffle=True, random_state=0)
    for train_index, test_index in kf.split(_X, _y):
        _clf.fit(_X[train_index], _y[train_index])
        y_pred = _clf.predict(_X[test_index])
        accuracies += [accuracy_score(_y[test_index], y_pred)]
    return np.array(accuracies)

acc = eval_classifier(GaussianNB(),
                      X, y)

print(f'Naive Bayes CV accuracy={np.mean(acc):.2f} {chr(177)}{np.std(acc):.3f}')
```

Naive Bayes CV accuracy=0.86 ±0.034

```
[60]: acc = eval_classifier(LinearSVC(dual=False, tol=10),
                             X, y)
print(f'Linear SVC CV accuracy={np.mean(acc):.2f} {chr(177)}{np.std(acc):.3f}')
```

Linear SVC CV accuracy=0.45 ±0.002

```
[61]: acc = eval_classifier(MLPClassifier(hidden_layer_sizes=(50,100),max_iter=500),
                             X, y)
print(f'Neural Network CV accuracy={np.mean(acc):.2f} {chr(177)}{np.std(acc):.3f}')
```

Neural Network CV accuracy=0.82 ±0.052

```
[62]: acc = eval_classifier(DecisionTreeClassifier(),
                             X, y)
print(f'Decision Tree CV accuracy={np.mean(acc):.2f} {chr(177)}{np.std(acc):.
↳3f}')

```

Decision Tree CV accuracy=0.79 ±0.046

```
[63]: acc = eval_classifier(RandomForestClassifier(n_estimators=200, max_depth=5,
↳random_state=0, n_jobs=4),
                             X, y)
print(f'Random Forest CV accuracy={np.mean(acc):.2f} {chr(177)}{np.std(acc):.
↳3f}')

```

Random Forest CV accuracy=0.88 ±0.039

0.2 Problem 2

```
[40]: mlp_weak = MLPClassifier(hidden_layer_sizes=(3, 3), max_iter=30, tol=1e-1)
dtc_weak = DecisionTreeClassifier(max_depth=5, max_features=5)

```

```
[127]: gnb_ensemble = [GaussianNB() for _ in range(100)]
svc_ensemble = [SVC(kernel='linear', probability=True, max_iter=1000) for _ in
↳range(100)]
mlp_ensemble = [MLPClassifier(hidden_layer_sizes=(3, 3), max_iter=30, tol=1e-1)
↳for _ in range(100)]
dtc_ensemble = [DecisionTreeClassifier(max_depth=5, max_features=5) for _ in
↳range(100)]

def eval_first_classifier(ensemble, X, y):
    first_classifier = ensemble[0]
    first_classifier.fit(X, y)
    score = first_classifier.score(X, y)
    return score

gnb_first_score = eval_first_classifier(gnb_ensemble, X, y)
svc_first_score = eval_first_classifier(svc_ensemble, X, y)
mlp_first_score = eval_first_classifier(mlp_ensemble, X, y)
dtc_first_score = eval_first_classifier(dtc_ensemble, X, y)

ensemble_results = {
    "Classifier": ["GaussianNB", "LinearSVC", "MLPClassifier",
↳"DecisionTreeClassifier"],
    "First Classifier Score": [gnb_first_score, svc_first_score,
↳mlp_first_score, dtc_first_score]
}

```

```
[128]: ensemble_results

```

```
[128]: {'Classifier': ['GaussianNB',
    'LinearSVC',
    'MLPClassifier',
    'DecisionTreeClassifier'],
    'First Classifier Score': [0.8616557734204793,
    0.5729847494553377,
    0.4466230936819172,
    0.8812636165577342]}
```

0.3 Problem 3

```
[112]: import random
from sklearn.utils import resample

def ensemble_fit(ensemble, X, y):
    for clf in ensemble:
        # Generate a bootstrap sample
        X_sample, y_sample = resample(X, y, replace=True, n_samples=int(0.8 *
↪len(X)))
        # Fit the classifier on the bootstrap sample
        clf.fit(X_sample, y_sample)

def report_first_classifier_performance(ensemble, X, y):
    first_clf = ensemble[0]
    y_pred = first_clf.predict(X)
    accuracy = accuracy_score(y, y_pred)
    print(f"Accuracy of the first classifier in the ensemble: {accuracy:.2f}")

ensemble_fit(gnb_ensemble, X, y)

print("GaussianNB Ensemble:")
report_first_classifier_performance(gnb_ensemble, X, y)
```

GaussianNB Ensemble:
Accuracy of the first classifier in the ensemble: 0.84

```
[113]: ensemble_fit(svc_ensemble, X, y)

print("\nSVC Ensemble:")
report_first_classifier_performance(svc_ensemble, X, y)
```

SVC Ensemble:
Accuracy of the first classifier in the ensemble: 0.61

```
[79]: ensemble_fit(mlp_ensemble, X, y)

print("\nMLPClassifier Ensemble:")
report_first_classifier_performance(mlp_ensemble, X, y)
```

MLPClassifier Ensemble:
Accuracy of the first classifier in the ensemble: 0.55

```
[80]: ensemble_fit(dtc_ensemble, X, y)

print("\nDecisionTreeClassifier Ensemble:")
report_first_classifier_performance(dtc_ensemble, X, y)
```

DecisionTreeClassifier Ensemble:
Accuracy of the first classifier in the ensemble: 0.85

0.4 Problem 4

```
[118]: ensembles = {
    "GaussianNB Ensemble": gnb_ensemble,
    "SVC Ensemble": svc_ensemble,
    "MLPClassifier Ensemble": mlp_ensemble,
    "DecisionTreeClassifier Ensemble": dtc_ensemble
}
for name, ensemble in ensembles.items():
    print(f"\nFitting {name}...")
    ensemble_fit(ensemble, X, y)
    print(f"\n{name} Performance:")
    report_first_classifier_performance(ensemble, X, y)

def ensemble_predict(ensemble, X):
    probas = np.zeros((X.shape[0], 2))
    for clf in ensemble:
        probas += clf.predict_proba(X)
    avg_probas = probas / len(ensemble)
    final_predictions = np.argmax(avg_probas, axis=1)
    return final_predictions
```

Fitting GaussianNB Ensemble...

GaussianNB Ensemble Performance:
Accuracy of the first classifier in the ensemble: 0.86

Fitting SVC Ensemble...

SVC Ensemble Performance:

Accuracy of the first classifier in the ensemble: 0.59

Fitting MLPClassifier Ensemble...

MLPClassifier Ensemble Performance:

Accuracy of the first classifier in the ensemble: 0.50

Fitting DecisionTreeClassifier Ensemble...

DecisionTreeClassifier Ensemble Performance:

Accuracy of the first classifier in the ensemble: 0.87

```
[119]: sample_input = X[:1]

gnb_predictions = ensemble_predict(gnb_ensemble, sample_input)
print("GaussianNB Ensemble Prediction:", gnb_predictions)
```

GaussianNB Ensemble Prediction: [0]

```
[120]: sample_input = X[:1]

gnb_predictions = ensemble_predict(gnb_ensemble, sample_input)
svc_predictions = ensemble_predict(svc_ensemble, sample_input)
mlp_predictions = ensemble_predict(mlp_ensemble, sample_input)
dtc_predictions = ensemble_predict(dtc_ensemble, sample_input)

print("GaussianNB Ensemble Prediction:", gnb_predictions)
print("SVC Ensemble Prediction:", svc_predictions)
print("MLPClassifier Ensemble Prediction:", mlp_predictions)
print("DecisionTreeClassifier Ensemble Prediction:", dtc_predictions)
```

GaussianNB Ensemble Prediction: [0]

SVC Ensemble Prediction: [1]

MLPClassifier Ensemble Prediction: [1]

DecisionTreeClassifier Ensemble Prediction: [0]

0.5 Problem 5

```
[143]: def ensemble_fit_with_subsample(ensemble, X, y, subsample_ratio):
        for clf in ensemble:
            while True:
                X_sample, y_sample = resample(X, y, replace=True,
                ↪n_samples=int(subsample_ratio * len(X)))
                if len(np.unique(y_sample)) > 1:
                    break
            clf.fit(X_sample, y_sample)
```



```

def evaluate_ensemble(ensemble, X, y, subsample_ratio, cv_folds=10):
    kf = StratifiedKFold(n_splits=cv_folds, shuffle=True, random_state=1)
    scores = []

    for train_index, test_index in kf.split(X, y):
        X_train, X_test = X[train_index], X[test_index]
        y_train, y_test = y[train_index], y[test_index]

        ensemble_fit_with_subsample(ensemble, X_train, y_train, subsample_ratio)

        predictions = ensemble_predict(ensemble, X_test)

        accuracy = accuracy_score(y_test, predictions)
        scores.append(accuracy)

    return np.array(scores)

for name, ensemble in ensembles.items():
    scores = evaluate_ensemble(ensemble, X, y, subsample_ratio=0.2, cv_folds=10)
    print(f"{name} 10-fold CV Mean Accuracy: {np.mean(scores):.4f} ± {np.
↪std(scores):.4f}")

```

GaussianNB Ensemble 10-fold CV Mean Accuracy: 0.8595 ± 0.0383

SVC Ensemble 10-fold CV Mean Accuracy: 0.5196 ± 0.0430

MLPClassifier Ensemble 10-fold CV Mean Accuracy: 0.5284 ± 0.0292

DecisionTreeClassifier Ensemble 10-fold CV Mean Accuracy: 0.8486 ± 0.0521

```

[126]: def evaluate_decision_tree(X, y, subsample_ratio, cv_folds=10):

    kf = StratifiedKFold(n_splits=cv_folds, shuffle=True, random_state=1)
    scores = []
    for train_index, test_index in kf.split(X, y):
        X_train, X_test = X[train_index], X[test_index]
        y_train, y_test = y[train_index], y[test_index]

        X_sample, y_sample = resample(X_train, y_train, replace=True,
↪n_samples=int(subsample_ratio * len(X_train)))

        clf = DecisionTreeClassifier(max_depth=5, max_features=5)
        clf.fit(X_sample, y_sample)

        y_pred = clf.predict(X_test)

        accuracy = accuracy_score(y_test, y_pred)
        scores.append(accuracy)

    return np.array(scores)

```

```
dt_scores_005 = evaluate_decision_tree(X, y, subsample_ratio=0.05, cv_folds=10)
print(f"Regular Decision Tree 10-fold CV Mean Accuracy with subsample ratio 0.05: {np.mean(dt_scores_005):.4f} ± {np.std(dt_scores_005):.4f}")
```

Regular Decision Tree 10-fold CV Mean Accuracy with subsample ratio 0.05: 0.7614 ± 0.0371

0.6 Problem 6

```
[149]: def create_ensembles():
    gnb_ensemble = [GaussianNB() for _ in range(10)]
    svc_ensemble = [SVC(kernel='linear', probability=True, max_iter=30) for _ in range(10)]
    mlp_ensemble = [MLPClassifier(hidden_layer_sizes=(3, 3), max_iter=30, tol=1e-1) for _ in range(10)]
    dtc_ensemble = [DecisionTreeClassifier(max_depth=5, max_features=5) for _ in range(10)]
    return {
        "GaussianNB Ensemble": gnb_ensemble,
        "SVC Ensemble": svc_ensemble,
        "MLPClassifier Ensemble": mlp_ensemble,
        "DecisionTreeClassifier Ensemble": dtc_ensemble
    }

def evaluate_single_classifier(clf, X, y, subsample_ratio, cv_folds=10):
    kf = StratifiedKFold(n_splits=cv_folds, shuffle=True, random_state=1)
    scores = []

    for train_index, test_index in kf.split(X, y):
        X_train, X_test = X[train_index], X[test_index]
        y_train, y_test = y[train_index], y[test_index]

        while True:
            X_sample, y_sample = resample(X_train, y_train, replace=True, n_samples=int(subsample_ratio * len(X_train)))
            if len(np.unique(y_sample)) > 1:
                break

        clf.fit(X_sample, y_sample)

        y_pred = clf.predict(X_test)
        accuracy = accuracy_score(y_test, y_pred)
        scores.append(accuracy)
```

```

    return np.array(scores)

subsample_ratios = [0.005, 0.01, 0.03, 0.05, 0.1, 0.2]

results = {
    "GaussianNB": {"ensemble": [], "regular": []},
    "SVC": {"ensemble": [], "regular": []},
    "MLPClassifier": {"ensemble": [], "regular": []},
    "DecisionTreeClassifier": {"ensemble": [], "regular": []}
}

# Evaluate ensembles and regular classifiers for each subsample ratio
for subsample_ratio in subsample_ratios:
    print(f"\nEvaluating with subsample ratio of {subsample_ratio}:")

    # Create fresh ensembles for each subsample ratio
    ensembles = create_ensembles()

    for name, ensemble in ensembles.items():
        scores = evaluate_ensemble(ensemble, X, y, subsample_ratio, cv_folds=10)
        classifier_type = name.split()[0]
        results[classifier_type]["ensemble"].append((subsample_ratio, np.
↳mean(scores), np.std(scores)))
        print(f"{name} 10-fold CV Mean Accuracy: {np.mean(scores):.4f} ± {np.
↳std(scores):.4f}")

    # Evaluate regular classifiers
    regular_classifiers = {
        "GaussianNB": GaussianNB(),
        "SVC": SVC(kernel='linear', probability=True, max_iter=30),
        "MLPClassifier": MLPClassifier(hidden_layer_sizes=(3, 3), max_iter=30,
↳tol=1e-1),
        "DecisionTreeClassifier": DecisionTreeClassifier(max_depth=5,
↳max_features=5)
    }

    for name, clf in regular_classifiers.items():
        scores = evaluate_single_classifier(clf, X, y, subsample_ratio,
↳cv_folds=10)
        classifier_type = name
        results[classifier_type]["regular"].append((subsample_ratio, np.
↳mean(scores), np.std(scores)))
        print(f"{name} with subsample ratio {subsample_ratio} 10-fold CV Mean
↳Accuracy: {np.mean(scores):.4f} ± {np.std(scores):.4f}")

```

Evaluating with subsample ratio of 0.005:

GaussianNB Ensemble 10-fold CV Mean Accuracy: 0.6470 ± 0.1210
SVC Ensemble 10-fold CV Mean Accuracy: 0.5044 ± 0.0803
MLPClassifier Ensemble 10-fold CV Mean Accuracy: 0.4945 ± 0.0449
DecisionTreeClassifier Ensemble 10-fold CV Mean Accuracy: 0.7627 ± 0.0838
GaussianNB with subsample ratio 0.005 10-fold CV Mean Accuracy: 0.5325 ± 0.1012
SVC with subsample ratio 0.005 10-fold CV Mean Accuracy: 0.5623 ± 0.0785
MLPClassifier with subsample ratio 0.005 10-fold CV Mean Accuracy: 0.5263 ± 0.0654
DecisionTreeClassifier with subsample ratio 0.005 10-fold CV Mean Accuracy: 0.6395 ± 0.1228

Evaluating with subsample ratio of 0.01:

GaussianNB Ensemble 10-fold CV Mean Accuracy: 0.7995 ± 0.0599
SVC Ensemble 10-fold CV Mean Accuracy: 0.6263 ± 0.0774
MLPClassifier Ensemble 10-fold CV Mean Accuracy: 0.5350 ± 0.0742
DecisionTreeClassifier Ensemble 10-fold CV Mean Accuracy: 0.8203 ± 0.0363
GaussianNB with subsample ratio 0.01 10-fold CV Mean Accuracy: 0.6577 ± 0.1053
SVC with subsample ratio 0.01 10-fold CV Mean Accuracy: 0.5752 ± 0.1100
MLPClassifier with subsample ratio 0.01 10-fold CV Mean Accuracy: 0.5097 ± 0.0716
DecisionTreeClassifier with subsample ratio 0.01 10-fold CV Mean Accuracy: 0.6675 ± 0.1351

Evaluating with subsample ratio of 0.03:

GaussianNB Ensemble 10-fold CV Mean Accuracy: 0.8519 ± 0.0360
SVC Ensemble 10-fold CV Mean Accuracy: 0.5631 ± 0.0560
MLPClassifier Ensemble 10-fold CV Mean Accuracy: 0.4935 ± 0.0563
DecisionTreeClassifier Ensemble 10-fold CV Mean Accuracy: 0.8258 ± 0.0340
GaussianNB with subsample ratio 0.03 10-fold CV Mean Accuracy: 0.7310 ± 0.0567
SVC with subsample ratio 0.03 10-fold CV Mean Accuracy: 0.5697 ± 0.0434
MLPClassifier with subsample ratio 0.03 10-fold CV Mean Accuracy: 0.5240 ± 0.0628
DecisionTreeClassifier with subsample ratio 0.03 10-fold CV Mean Accuracy: 0.7082 ± 0.0784

Evaluating with subsample ratio of 0.05:

GaussianNB Ensemble 10-fold CV Mean Accuracy: 0.8453 ± 0.0400
SVC Ensemble 10-fold CV Mean Accuracy: 0.5152 ± 0.0766
MLPClassifier Ensemble 10-fold CV Mean Accuracy: 0.5098 ± 0.0741
DecisionTreeClassifier Ensemble 10-fold CV Mean Accuracy: 0.8290 ± 0.0364
GaussianNB with subsample ratio 0.05 10-fold CV Mean Accuracy: 0.7953 ± 0.0395
SVC with subsample ratio 0.05 10-fold CV Mean Accuracy: 0.5118 ± 0.0815
MLPClassifier with subsample ratio 0.05 10-fold CV Mean Accuracy: 0.4924 ± 0.0502
DecisionTreeClassifier with subsample ratio 0.05 10-fold CV Mean Accuracy: 0.7616 ± 0.0651

Evaluating with subsample ratio of 0.1:

GaussianNB Ensemble 10-fold CV Mean Accuracy: 0.8475 ± 0.0468

SVC Ensemble 10-fold CV Mean Accuracy: 0.5088 ± 0.0675

MLPClassifier Ensemble 10-fold CV Mean Accuracy: 0.5306 ± 0.0659

DecisionTreeClassifier Ensemble 10-fold CV Mean Accuracy: 0.8399 ± 0.0460

GaussianNB with subsample ratio 0.1 10-fold CV Mean Accuracy: 0.8312 ± 0.0438

SVC with subsample ratio 0.1 10-fold CV Mean Accuracy: 0.5339 ± 0.0906

MLPClassifier with subsample ratio 0.1 10-fold CV Mean Accuracy: 0.4989 ± 0.0524

DecisionTreeClassifier with subsample ratio 0.1 10-fold CV Mean Accuracy: 0.7408 ± 0.0508

Evaluating with subsample ratio of 0.2:

GaussianNB Ensemble 10-fold CV Mean Accuracy: 0.8584 ± 0.0462

SVC Ensemble 10-fold CV Mean Accuracy: 0.4879 ± 0.0688

MLPClassifier Ensemble 10-fold CV Mean Accuracy: 0.5011 ± 0.0556

DecisionTreeClassifier Ensemble 10-fold CV Mean Accuracy: 0.8584 ± 0.0458

GaussianNB with subsample ratio 0.2 10-fold CV Mean Accuracy: 0.8377 ± 0.0489

SVC with subsample ratio 0.2 10-fold CV Mean Accuracy: 0.5926 ± 0.0242

MLPClassifier with subsample ratio 0.2 10-fold CV Mean Accuracy: 0.5205 ± 0.0659

DecisionTreeClassifier with subsample ratio 0.2 10-fold CV Mean Accuracy: 0.7855 ± 0.0537

0.7 Problem 7

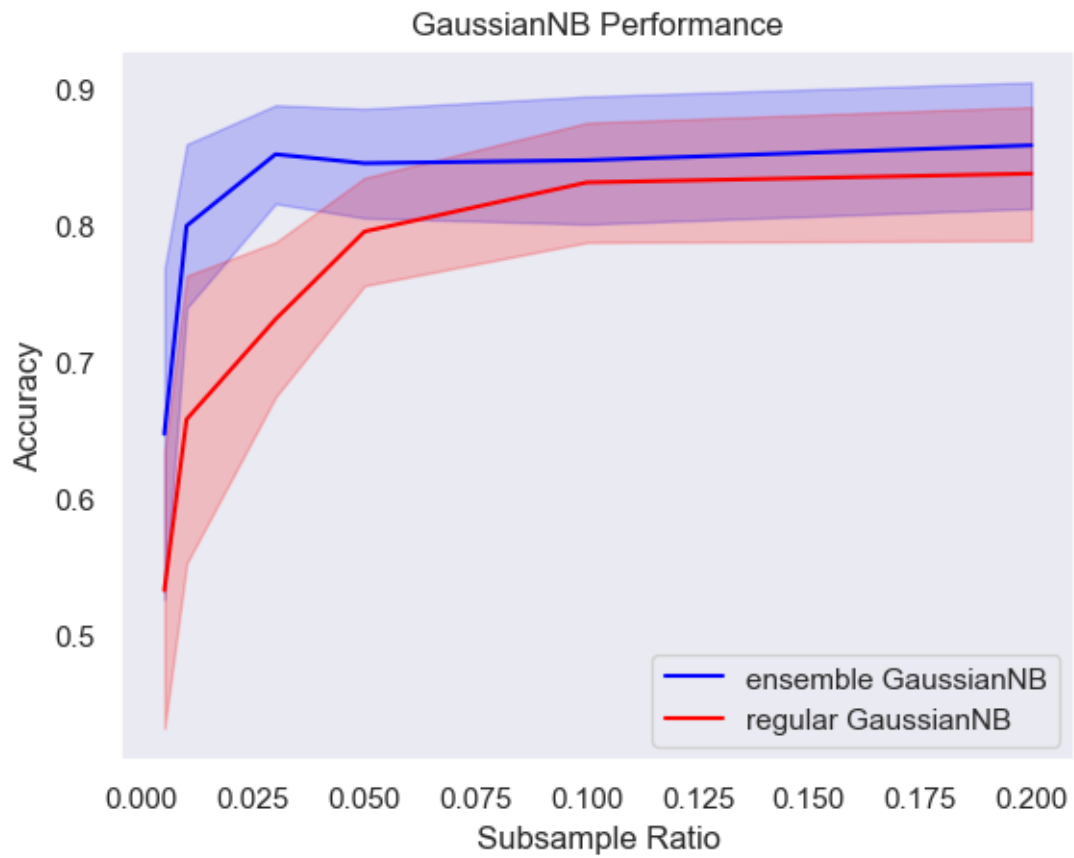
```
[152]: colors = {
        "ensemble": "blue",
        "regular": "red"
    }

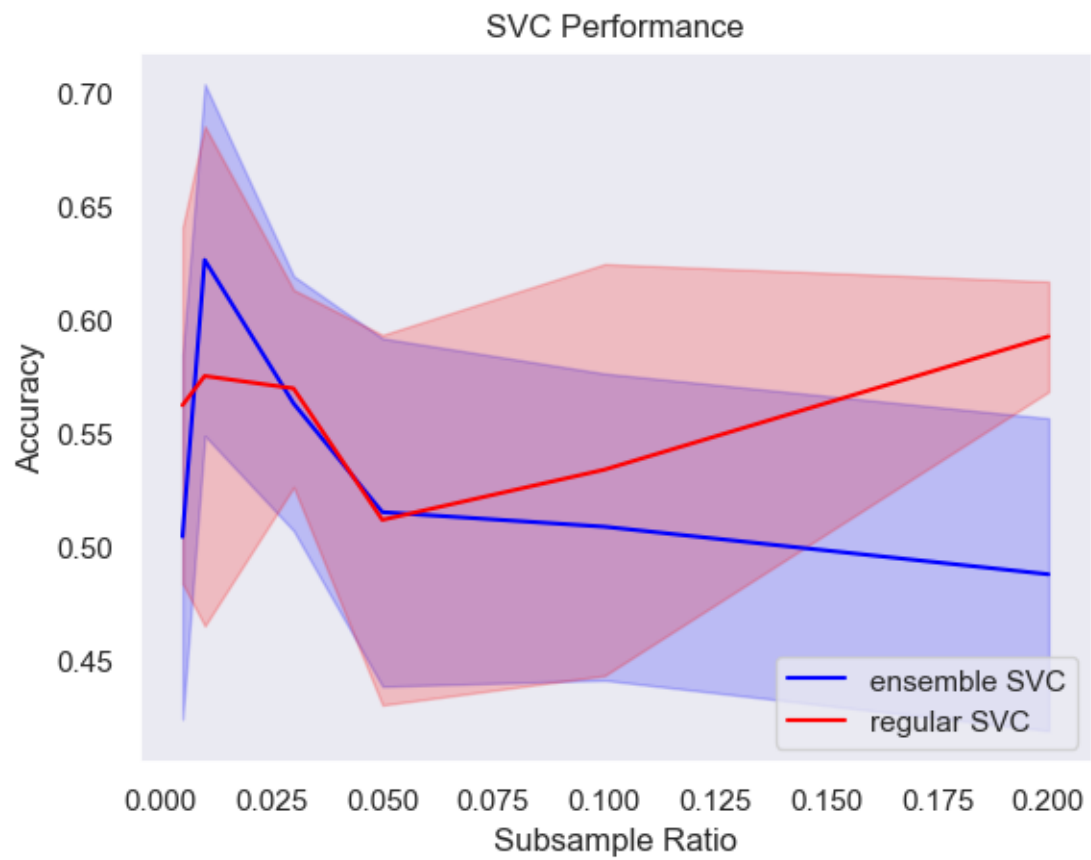
    for classifier_type, data in results.items():
        plt.figure()
        for label, color in colors.items():
            vals, accs, stdevs = zip(*data[label])

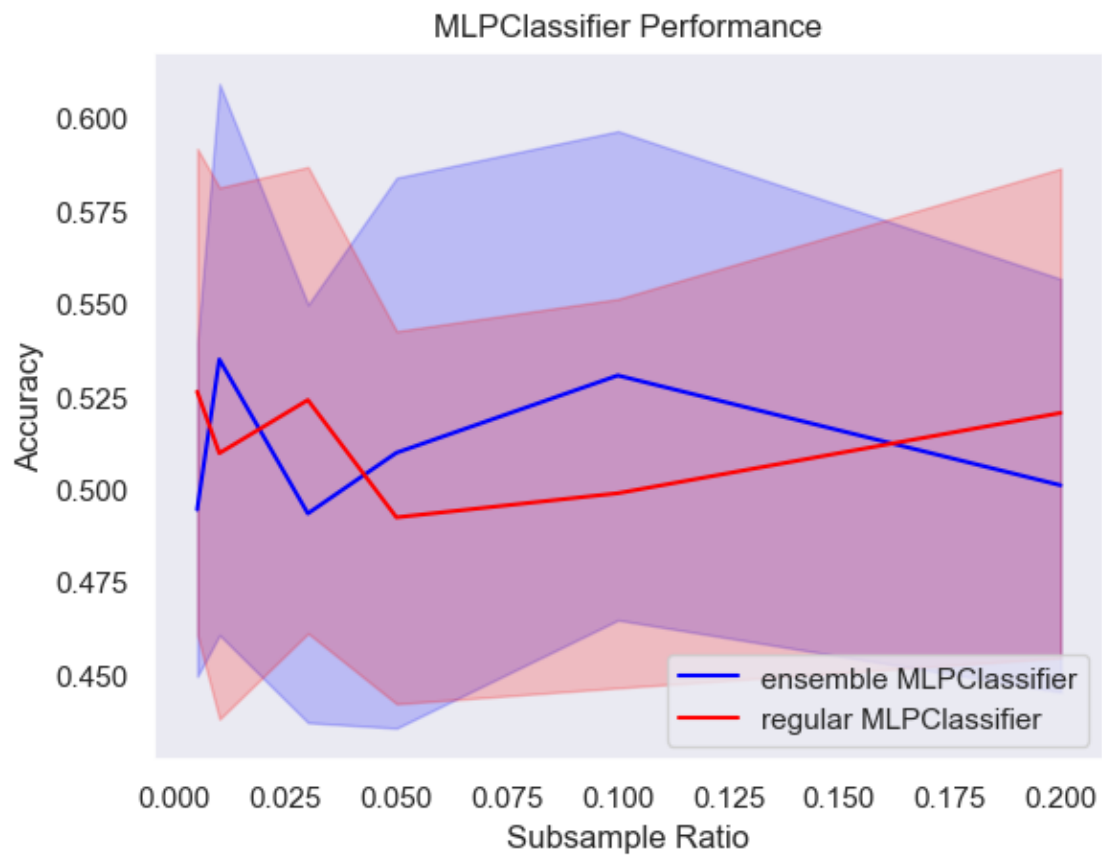
            vals, accs, stdevs = np.array(vals), np.array(accs), np.array(stdevs)


            plt.plot(vals, accs, label=f'{label} {classifier_type}', color=color)
            plt.fill_between(vals, accs - stdevs, accs + stdevs, color=color,
                             alpha=0.2)

        plt.xlabel('Subsample Ratio')
        plt.ylabel('Accuracy')
        plt.legend(loc='lower right')
        plt.grid()
        plt.title(f'{classifier_type} Performance')
        plt.show()
```









[]: